

# Performance issues

# Indexes

- Indexes are a common way to **enhance database performance**.
  - An index allows the database server to find and retrieve specific rows **much faster** than it could do without an index.
  - But indexes also **add overhead to the database system** as a whole, so they should be used sensibly

# Create index

- **CREATE INDEX** test1\_id\_index **ON** test1 (id);
  - **CREATE INDEX** test1\_id\_index **ON** test1 **USING** btree (id);
  - **CREATE INDEX** test1\_id\_index **ON** test1 [USING btree] (id) **WHERE** <condition>;
- ➔ Partial index

# Index types in PostgreSQL

- PostgreSQL provides several index types: B-tree, Hash, GiST, SP-GiST, GIN, BRIN
- Each index type uses a different algorithm that is best suited to different types of queries.

By default, the CREATE INDEX command creates **B-tree** indexes, which **fit the most common situations**

# Index types in PostgreSQL

- **B-Tree** (default)
  - handle equality and range queries on data that can be sorted into some ordering.
  - Operators:  $<$ ,  $\leq$ ,  $=$ ,  $\geq$ ,  $>$ , LIKE (col LIKE 'foo%' but not col LIKE '%bar')
  - Sorted output
- **Hash index**: can only handle simple equality comparisons
- **GiST index**: for several two-dimensional geometric data types,
  - not a single kind of index, but rather an infrastructure within which many different indexing strategies can be implemented
- **GIN index**
  - inverted indexes which can handle values that contain more than one key, arrays for example

# Index types in PostgreSQL

- **spgist index:** .....
- **Brin:** .....

# Multicolumn index

- **CREATE INDEX** test2\_mm\_idx **ON** test2 (major, minor);
- B-Tree
- GiST index
- GIN index


<https://www.postgresql.org/docs/10/sql-createindex.html>

<https://www.postgresql.org/docs/10/indexes.html>

# Examining index usage

- **EXPLAIN [ ANALYZE ] [ VERBOSE ] *statement***
  - **EXPLAIN *statement***: displays the execution plan that the PostgreSQL planner generates for the supplied statement.

Actually two numbers are shown: the **start-up cost** before the first row can be returned, and the **total cost** to return all the rows.
  - **VERBOSE** option: displays additional information regarding the plan (output column list, table and function names, ...)





# Examining index usage

- **EXPLAIN [ ANALYZE ] [ VERBOSE ] *statement***
  - **ANALYZE** option: causes the statement to be **actually executed**, not only planned, actual runtime statistics are added to the display
- **Important:** If you wish to use **EXPLAIN ANALYZE** on an INSERT, UPDATE, DELETE, CREATE TABLE AS, or EXECUTE statement without letting the command affect your data, use this approach:

**BEGIN;**

**EXPLAIN ANALYZE ...;**

**ROLLBACK;**

# View table indexes

- \d table\_name
- Ex.: \d customers

# Tips

- Select **fewer columns** to improve hash join performance
- Index the *independent* **where** predicates to improve hash join performance

# Tips

- Having a **WHERE / HAVING** clause in your queries does not necessarily mean that it is a bad query
- Only retrieve the data you need
  - remove unnecessary columns from **SELECT**
  - **Inner join** vs. **exists** (with subqueries)
  - Select **DISTINCT** : try to avoid if you can
  - **LIKE** operator: the **index isn't used** if the pattern **starts with % or \_**

- Limit your results : LIMIT, TOP
- Don't Make Queries More Complex Than They Need To Be
  - OR / IN / UNION ?
  - OR operator : index is not used except composite index → IN/UNION/OUTER JOIN
  - NOT operator: index is not used => avoid
  - AND vs BETWEEN
  - ANY / ALL: index not used => max , min ,...
  - Isolate columns in Condition :  $\text{age} + 7 < 20 \rightarrow \text{age} < 13$

- Limit your results : **LIMIT, TOP**

You can add the LIMIT or TOP clauses to your queries to set a maximum number of rows for the result set.

```
SELECT TOP 3 *  
FROM customers;
```

```
SELECT *  
FROM customers  
LIMIT 3;
```

- Don't Make Queries More Complex Than They Need To Be
  - OR / IN / UNION ?
  - OR operator : index is not used except composite index → IN/UNION/OUTER JOIN
    - Using a condition with IN or UNION:

```
SELECT * FROM orderlines
WHERE orderid = 1 OR orderid = 5000;
-- (first cost: 8 - total cost: 47).
Actual time = 50.82..50.83
```

```
SELECT * FROM orderlines
WHERE orderid IN (1,5000);
-- (0.29 - 30), actual time = 0.028..0.039
```

```
SELECT * FROM orderlines
WHERE orderid = 1
UNION
SELECT * FROM orderlines
WHERE orderid = 5000;
-- (30 - 31) - actual time: 0.053..0.056
```



- Don't Make Queries More Complex Than They Need To Be
  - To be careful not to unnecessarily use the UNION operation because you go through the same table multiple times → use a UNION in your query, the execution time will increase.
  - Alternatives to the UNION operation are: reformulating the query in such a way that all conditions are placed in one SELECT instruction, or using an OUTER JOIN instead of UNION.

```
SELECT P.* , o.quantity
FROM products p left join orderlines o ON(p.prod_id =
o.prod_id) -- (326 - 2076), ~500ms
WHERE o.orderlineid IS NULL; -- (326 - 2076), 162ms
```

```
SELECT * , 0
FROM products
WHERE prod_id not in (select prod_id from orderlines)
UNION
SELECT p.*, quantity
FROM products p join orderlines o ON(p.prod_id = o.prod_id);
-- (17 780 - 19 210) 864 ms
```

Explain analyze

```
SELECT * , 0  
FROM products  
WHERE prod_id not in (select prod_id from orderlines)  
UNION  
SELECT p.*, quantity  
FROM products p join orderlines o ON(p.prod_id = o.prod_id);
```

Output Explain Messages Notifications

#### QUERY PLAN

text



Unique (cost=17463.15..18933.52 rows=65350 width=270) (actual time=267.156..358.138 rows=26098 loops=1)

-> Sort (cost=17463.15..17626.52 rows=65350 width=270) (actual time=267.154..328.189 rows=60377 loops=1)

Sort Key: products.prod\_id, products.category, products.title, products.actor, products.price, products.special, products.common\_prod...

Sort Method: external merge Disk: 3864kB

-> Append (cost=1139.38..3969.49 rows=65350 width=270) (actual time=45.064..125.939 rows=60377 loops=1)

-> Seq Scan on products (cost=1139.38..1365.38 rows=5000 width=53) (actual time=45.063..48.695 rows=27 loops=1)

Filter: (NOT (hashed SubPlan 1))

### Explain analyze

```
SELECT P.* , o.quantity  
FROM products p left join orderlines o ON (p.prod_id = o.prod_id)  
WHERE o.orderlineid IS NULL;
```

Output Explain Messages Notifications

#### QUERY PLAN

text

Hash Right Join (cost=326.00..1472.99 rows=1 width=51) (actual time=58.881..59.436 rows=27 loops=1)

Hash Cond: (o.prod\_id = p.prod\_id)

Filter: (o.orderlineid IS NULL)

Rows Removed by Filter: 60350

-> Seq Scan on orderlines o (cost=0.00..988.50 rows=60350 width=10) (actual time=0.023..10.794 rows=60350 loops=1)

-> Hash (cost=201.00..201.00 rows=10000 width=49) (actual time=2.543..2.545 rows=10000 loops=1)

Buckets: 16384 Batches: 1 Memory Usage: 943kB

-> Seq Scan on products p (cost=0.00..201.00 rows=10000 width=49) (actual time=0.014..0.735 rows=10000 loops=1)

Planning Time: 0.275 ms

Execution Time: 59.729 ms

- **NOT operator**: index is not used => avoid

```
select * from customers  
where customerid != 5000;
```

```
select * from customers  
where customerid = 5000;
```

- **ANY / ALL**: index not used => max , min ,...

- Isolate columns in Condition :

**age + 7 < 20** → **age < 13**

- No Brute force
    - JOIN clause:
      - Order of tables => biggest table: placed last in join
      - No redundant conditions on joins
    - Having clause:
      - Used only if needed
      - Not to replace WHERE => WHERE help to limit the intermediate number of records
- ➔ Need smart indexing, smart using

# Other index types

- Geometric type :
  - <https://www.postgresql.org/docs/10/datatype-geometric.html>
  - <https://www.postgresql.org/docs/10/functions-geometry.html>
- GiST:  
<https://www.postgresql.org/docs/10/indexes-types.html>